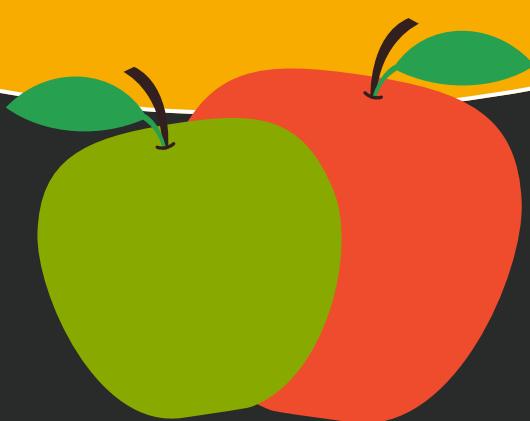


Classification d'images

Une comparaison
entre R et Python

Conférence Machine Learning



Clément Melina, Le Moan Delalande Riwal, Mathieu Anna

21-11-25

Sommaire

. 01

Introduction

. 02

Qu'est-ce qu'une API ?

. 03

Comment fonctionne Keras sur R et sur Python

. 04

Méthodologie

. 05

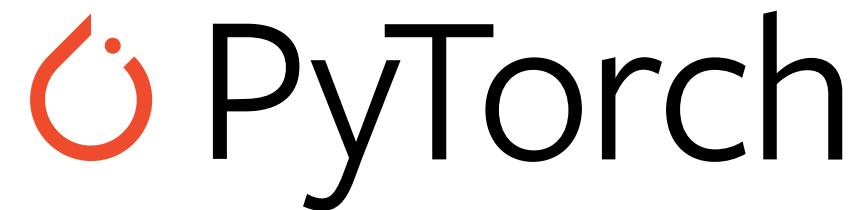
Comparaison des méthodes

. 06

Conclusion

Introduction Plusieurs approches pour le deep learning

Le deep learning peut être abordé à travers plusieurs outils et langages

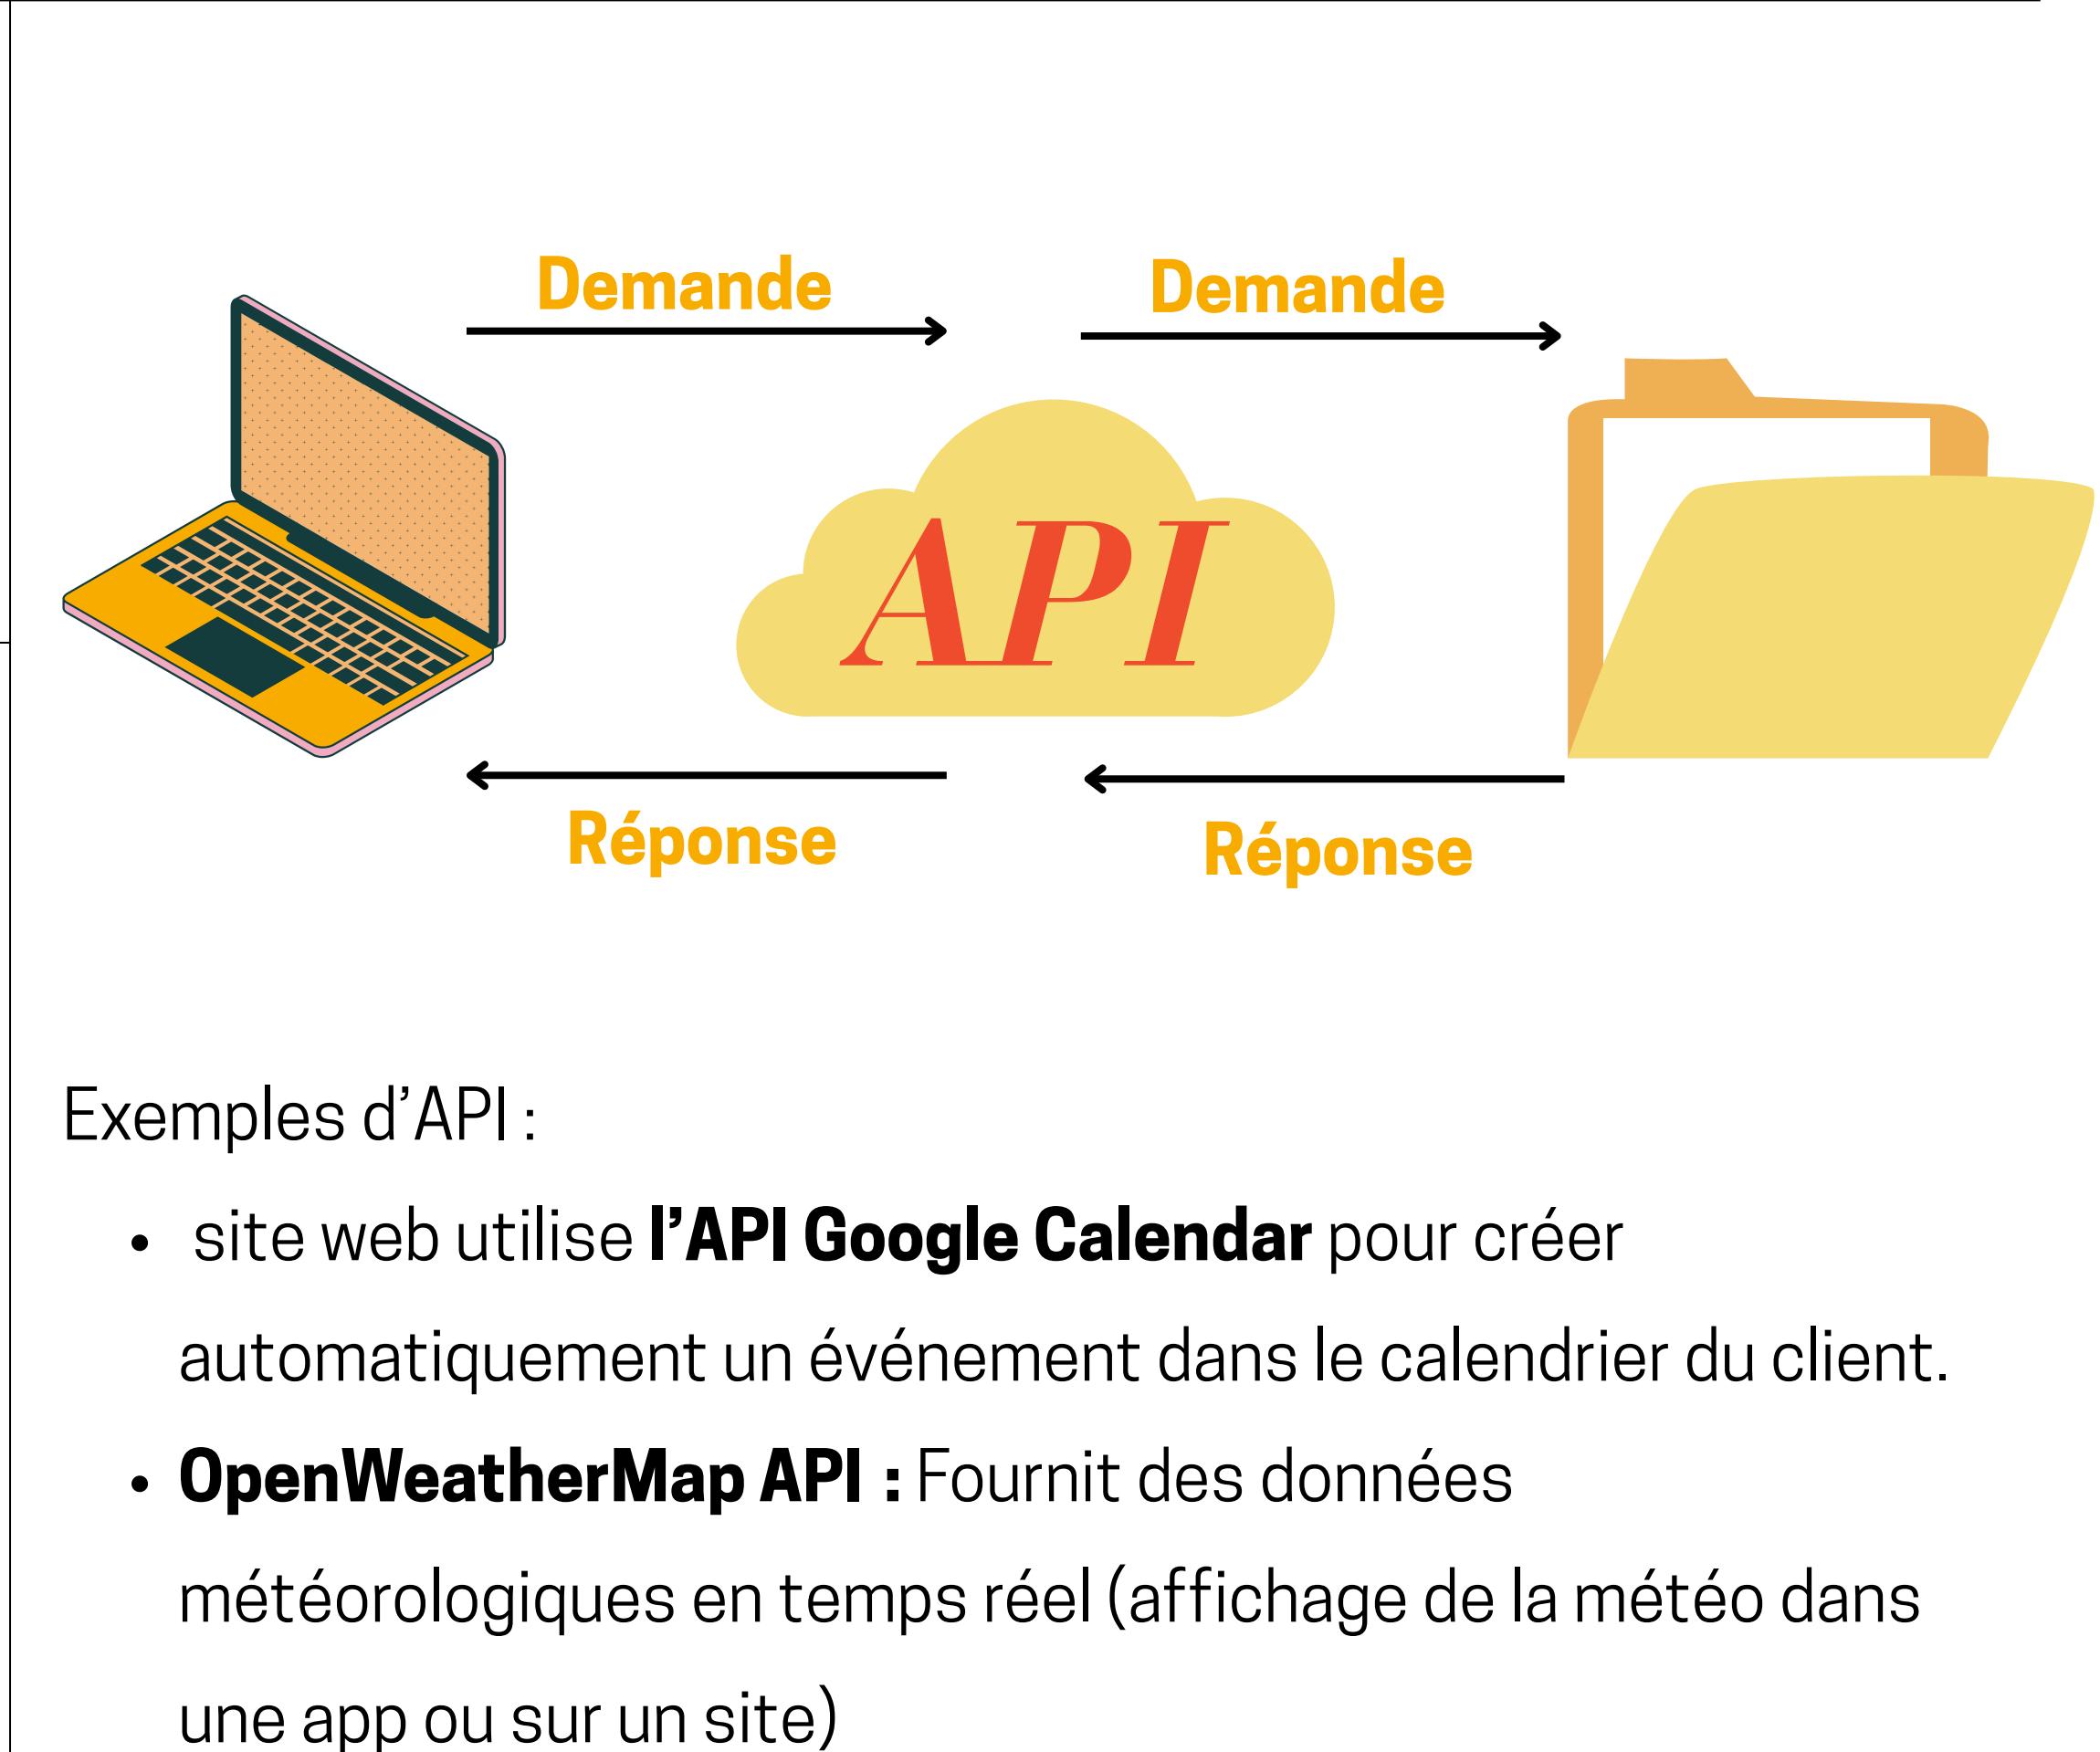


Comment un même projet de deep learning peut-il être abordé différemment selon l'outil utilisé ? Quels sont les avantages et limites de R Keras, Python Keras et PyTorch pour construire, entraîner et exploiter un modèle ?

Qu'est ce qu'une API

Application Programming Interface

- Connecte 2 logiciels/services ensemble
- Permet échange de données et de fonctionnalités



Keras



Objectifs :

Orienté humain

Vise à simplifier la façon de coder

Framework :

Permet d'utiliser plusieurs Framework de la même façon : JAX, TensorFlow, et PyTorch.

Utilisation :

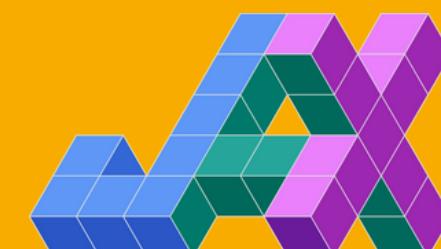
2 modes d'utilisation principaux :

- Séquentiel (Empilement de couches, CNN)
- Fonctionnel (Modèles plus complexes, non linéaires)



PyTorch

TensorFlow



Sur R

Package R permettant de faire le lien entre python et R

```
reticulate::install_python(version = "3.9")
```

Permet d'installer une version de python directement depuis R

Reticulate permet de traduire les objets R en objets python



K Keras

```
library(keras)
```

```
library(keras3)
```

interface R vers l'API Keras Python

Facilite la création de CNN dans une environnement R

Similaire à la structure Keras python

2 version du package coexistent

TensorFlow

```
library(TensorFlow)
```

Base pour l'utilisation du package keras

Permet la manipulation des Tenseurs

Nos données : classification d'image

Source :

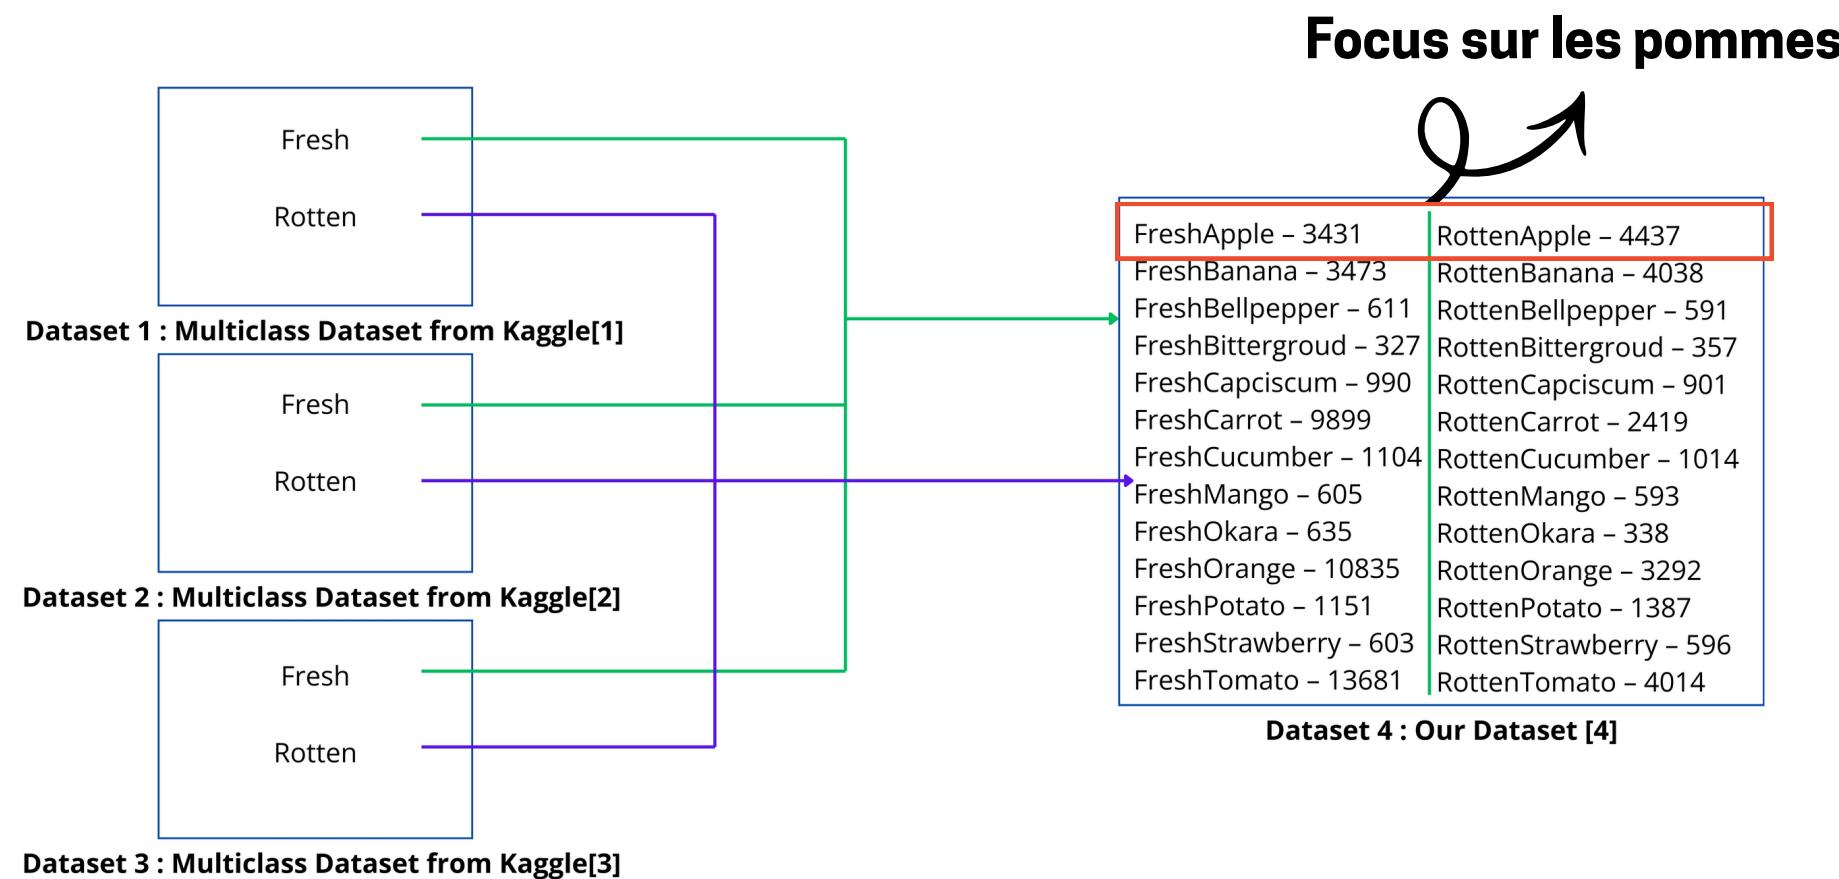


Food Freshness Dataset

Standardized Fruit and Vegetable Images for Quality Detection Tasks

kaggle.com

Le jeu de données : issu de la compilation de trois jeux de données de classification fruits frais/pourris



Composition

- > 70 000 images
 - Fruits et légumes frais
 - Fruits et légumes pourris
- 13 catégories de fruits et légumes
- taille : 6.4 GB

Fresh apple



Rotten apple



Obtention des données

- Preprocessing et data augmentation
- Standardisation en 128x128 pixels et transformations
 - Conversion en niveaux de gris
 - détection des contours
 - réglage de la luminosité

Méthodologie



Keras



TensorFlow

0) Préparation des fichiers

```

classes <- c("Mures", "Pourris")
train_dir <- file.path(base_dir, "train")
test_dir <- file.path(base_dir, "test")

dir.create(train_dir, recursive = TRUE, showWarnings = FALSE)
dir.create(test_dir, recursive = TRUE, showWarnings = FALSE)

for (class in classes) {
  cat("Traitement de la classe:", class, "\n")
  source_path <- file.path(base_dir, class)
  train_class_path <- file.path(train_dir, class)
  test_class_path <- file.path(test_dir, class)

  dir.create(train_class_path, showWarnings = FALSE)
  dir.create(test_class_path, showWarnings = FALSE)

  all_files <- list.files(source_path, full.names = TRUE)
  num_files <- length(all_files)

  if (num_files == 0) {
    cat(" !!! Aucun fichier trouvé dans", source_path, "!!!\n")
    next
  }

  train_indices <- sample(seq_len(num_files), size = floor(split_ratio * num_files))
  train_files <- all_files[train_indices]
  test_files <- all_files[-train_indices]

  file.copy(from = train_files, to = file.path(train_class_path, basename(train_files)))
  file.copy(from = test_files, to = file.path(test_class_path, basename(test_files)))

  cat(" →", length(train_files), "fichiers déplacés vers 'train/'", class, "\n", sep="")
  cat(" →", length(test_files), "fichiers déplacés vers 'test/'", class, "\n", sep="")
}

```

1) Préparation des données

```

image_size <- c(128, 128)

train_datagen <- image_data_generator(
  rescale = 1/255,
  shear_range = 0.2,
  zoom_range = 0.2,
  horizontal_flip = TRUE
)

test_datagen <- image_data_generator(
  rescale = 1/255
)

train_generator <- flow_images_from_directory(
  train_dir,
  generator = train_datagen,
  target_size = image_size,
  batch_size = 100,
  class_mode = 'binary'
)

test_generator <- flow_images_from_directory(
  test_dir,
  generator = test_datagen,
  target_size = image_size,
  batch_size = 100,
  class_mode = 'binary'
)

```

2) Construction du modèle

```

model <- keras_model_sequential() %>%
  layer_conv_2d(filters = 16, kernel_size = c(3, 3),
                activation = 'relu') %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 32, kernel_size = c(3, 3),
                activation = 'relu') %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_flatten() %>%
  layer_dense(units = 64, activation = 'relu') %>%
  layer_dense(units = 1, activation = 'sigmoid')

```

Méthodologie

③ Compilation et entraînement

```
model %>% compile(
  optimizer = optimizer_adam(),
  loss = 'binary_crossentropy',
  metrics = c('accuracy')
)
EPOCHS ← 10
options(keras.view_metrics = FALSE)

history ← model %>% fit(
  train_generator,
  steps_per_epoch = ceiling(train_generator$n / train_generator$batch_size),
  epochs = EPOCHS,
  validation_data = test_generator,
  validation_steps = ceiling(test_generator$n / test_generator$batch_size),
  view_metrics = FALSE |
)
```



④ Validation

```
evaluation ← model %>% evaluate(
  test_generator,
  steps = ceiling(test_generator$n / test_generator$batch_size)
)
|
cat('Test loss:', evaluation$loss, '\n')
cat('Test accuracy:', evaluation$accuracy, '\n')
```

Méthodologie



Keras



TensorFlow

① Préparation des données

```
train_ds = tf.keras.utils.image_dataset_from_directory(
    DATA_DIR,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(128, 128),
    batch_size=32,
    label_mode="binary"
)

val_ds = tf.keras.utils.image_dataset_from_directory(
    DATA_DIR,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(128, 128),
    batch_size=32,
    label_mode="binary"
)

train_ds = train_ds.prefetch(tf.data.AUTOTUNE)
val_ds = val_ds.prefetch(tf.data.AUTOTUNE)
```

② Construction du modèle

```
data_augmentation = keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomZoom(0.2),
    layers.RandomRotation(0.1),
])

model = Sequential([
    keras.Input(shape=(128, 128, 3)),
    data_augmentation,
    layers.Rescaling(1./255),
    layers.Conv2D(16, 3, activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])
```

③ Compilation et entraînement

```
model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=10
)
```

Méthodologie



① Préparation des données

```

transform = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ToTensor()
])

full_dataset = datasets.ImageFolder(root=path, transform=transform)

val_size = int(0.2 * len(full_dataset))
train_size = len(full_dataset) - val_size
train_dataset, val_dataset = random_split(
    full_dataset,
    [train_size, val_size],
    generator=torch.Generator().manual_seed(123)
)

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)

```

② Construction du modèle

```

class AppleCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 16, kernel_size=3)
        self.conv2 = nn.Conv2d(16, 32, kernel_size=3)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(32 * 30 * 30, 64)
        self.fc2 = nn.Linear(64, 1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, start_dim=1)
        x = F.relu(self.fc1(x))
        x = torch.sigmoid(self.fc2(x))
        return x

model = AppleCNN().to(device)
criterion = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)

train_acc_list, val_acc_list = [], []
train_loss_list, val_loss_list = [], []

```

Méthodologie

③ Compilation et entrainement

```

for epoch in range(EPOCHS):
    # ----- TRAIN -----
    model.train()
    running_loss = 0.0
    running_correct = 0
    total_train = 0

    for images, labels in train_loader:
        images = images.to(device)
        labels = labels.float().view(-1, 1).to(device) #

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * images.size(0)
        preds = (outputs >= 0.5).float()
        running_correct += (preds == labels).sum().item()
        total_train += labels.size(0)

    epoch_loss = running_loss / total_train
    epoch_acc = running_correct / total_train

```



④ Validation

```

# ----- VALIDATION -----
model.eval()
val_loss = 0.0
val_correct = 0
total_val = 0

with torch.no_grad():
    for images, labels in val_loader:
        images = images.to(device)
        labels = labels.float().view(-1, 1).to(device)

        outputs = model(images)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * images.size(0)
        preds = (outputs >= 0.5).float()
        val_correct += (preds == labels).sum().item()
        total_val += labels.size(0)

val_loss /= total_val
val_acc = val_correct / total_val

train_loss_list.append(epoch_loss)
train_acc_list.append(epoch_acc)
val_loss_list.append(val_loss)
val_acc_list.append(val_acc)

```

Gestion des données



Données locales uniquement (pas de chargement direct depuis le web)

Split manuel en dossiers train/ et test/
=> double stockage des données

Durée dépendante de la quantité de données :

- < 1 min | 1 Go
- ≈ 7 min | 2 Go
- ≈ 15 min | 3.5 Go

Chargement avec `image_data_generator()`

Pipeline simple mais moins flexible



Charge facilement :



Compatible datasets distants :

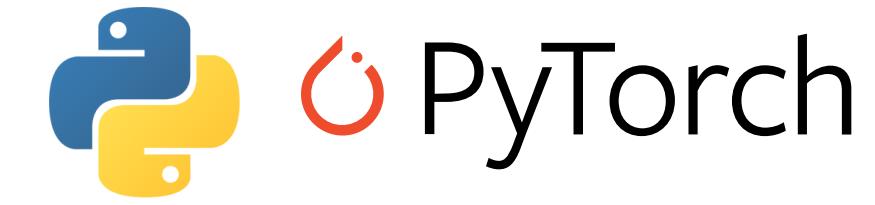


Chargement très automatisé :

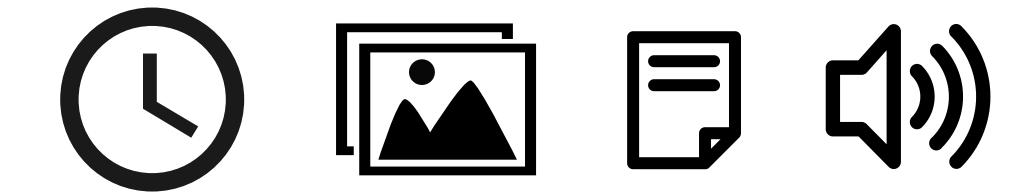
`tf.keras.utils.image_dataset_from_directory()`

Split intégré :

`validation_split=0.2,`



Charge facilement :



Compatible datasets distants :



Système de Dataset et de DataLoader :

`datasets.ImageFolder()`
`DataLoader()`

Split facilité :

`torch.utils.data.random_split()`

Transformations de données avancées

Construction du modèle



La construction du modèle via keras_model_sequential() avant l'ajout des couches avec des fonctions layer_...()

```
model <- keras_model_sequential()
model %>%
    ## layers....
```

Ensuite compilation du modèle :

```
model %>% compile(
  optimizer = 'rmsprop',
  loss = 'categorical_crossentropy',
  metrics = c('accuracy')
)
```

Et formation du modèle :

```
model %>% fit(x_train, y_train, epochs = 20, batch_size = 128)
```

Facile, rapide et intuitif

Différences uniquement liées au langage de programmation



La construction du modèle via Sequential() avec une facilité d'ajouter des couches

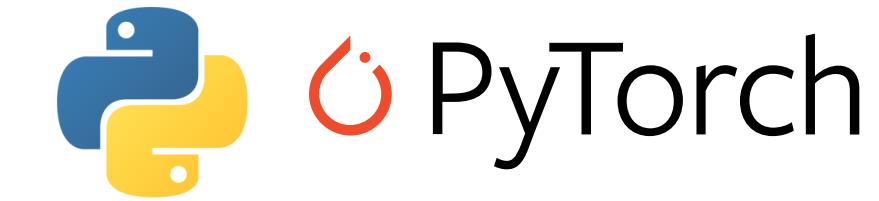
```
model = Sequential([
    ## layers....
```

Ensuite compilation du modèle :

```
model.compile(
  optimizer='adam',
  loss='binary_crossentropy',
  metrics=['accuracy']
)
```

Et formation du modèle :

```
history = model.fit(
  train_ds,
  validation_data=val_ds,
  epochs=10
)
```



Les modèles sont définis sous forme de classes (nn.Module)

```
class AppleCNN(nn.Module):
```

La construction de la structure du réseau :

```
def __init__(self):
  super().__init__()
  self.conv1 = nn.Conv2d(3, 32, kernel_size=3)

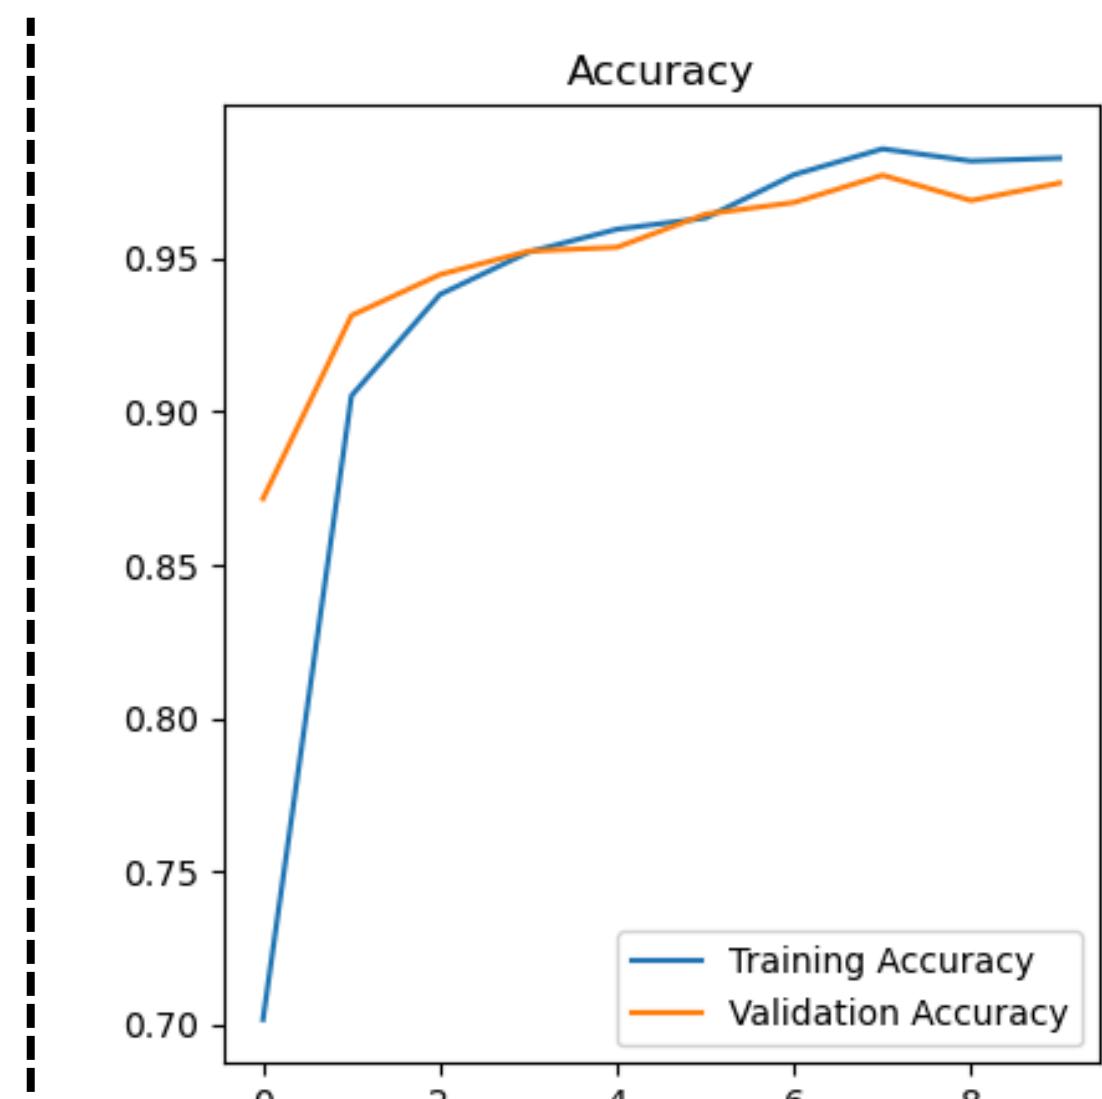
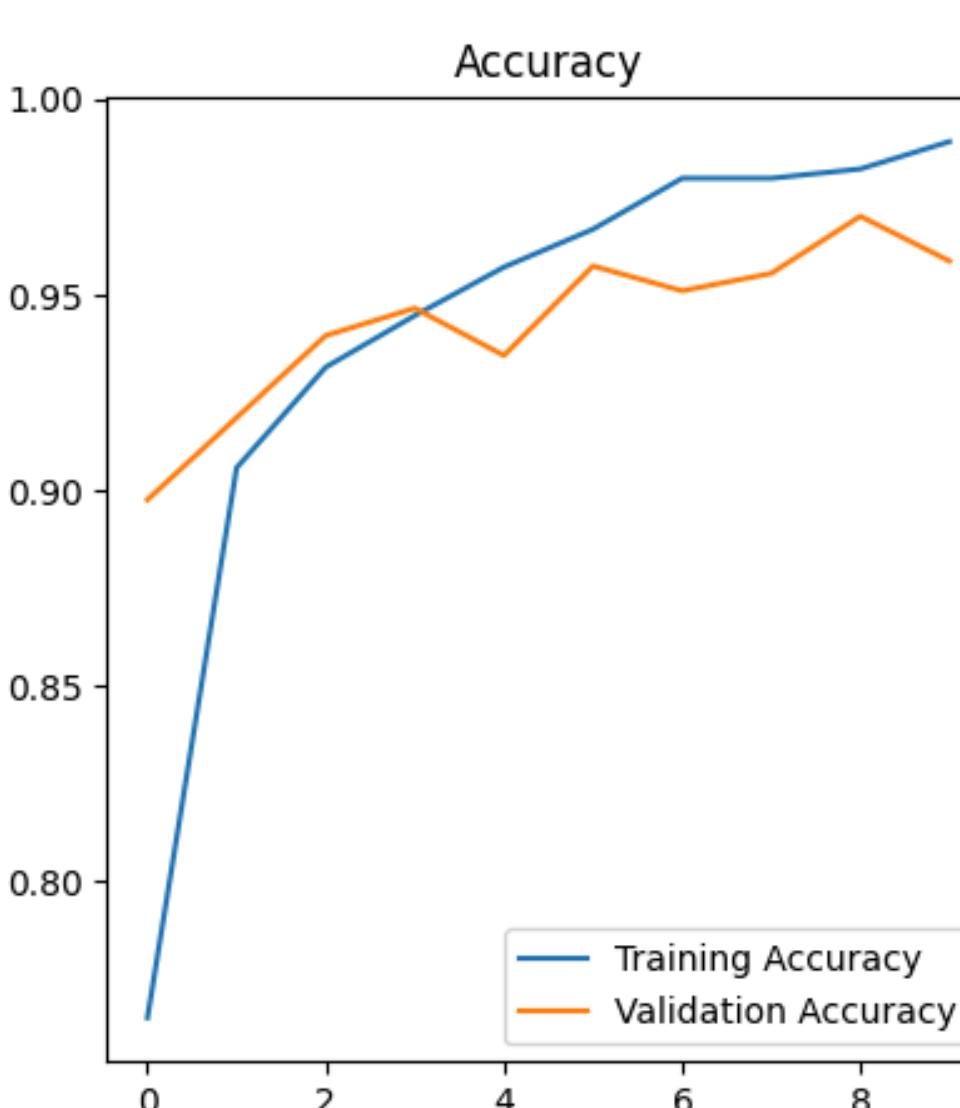
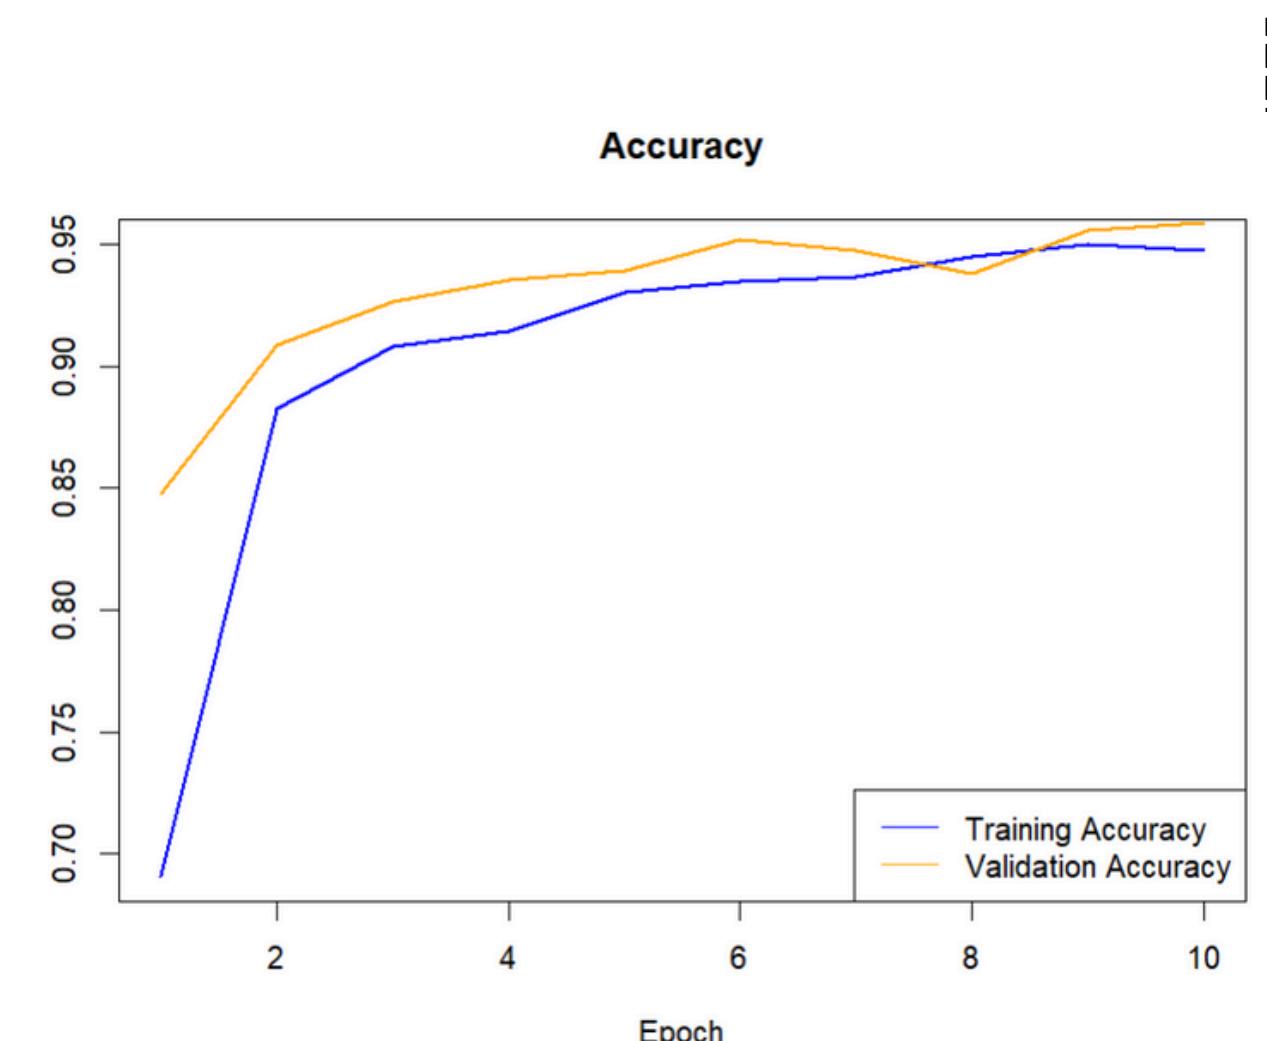
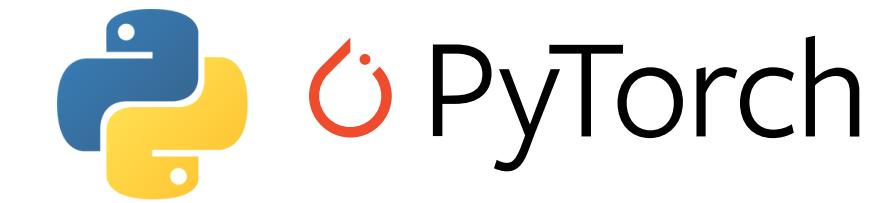
  ...
  def forward(self, x):
```

Logique de compréhension des tenseurs

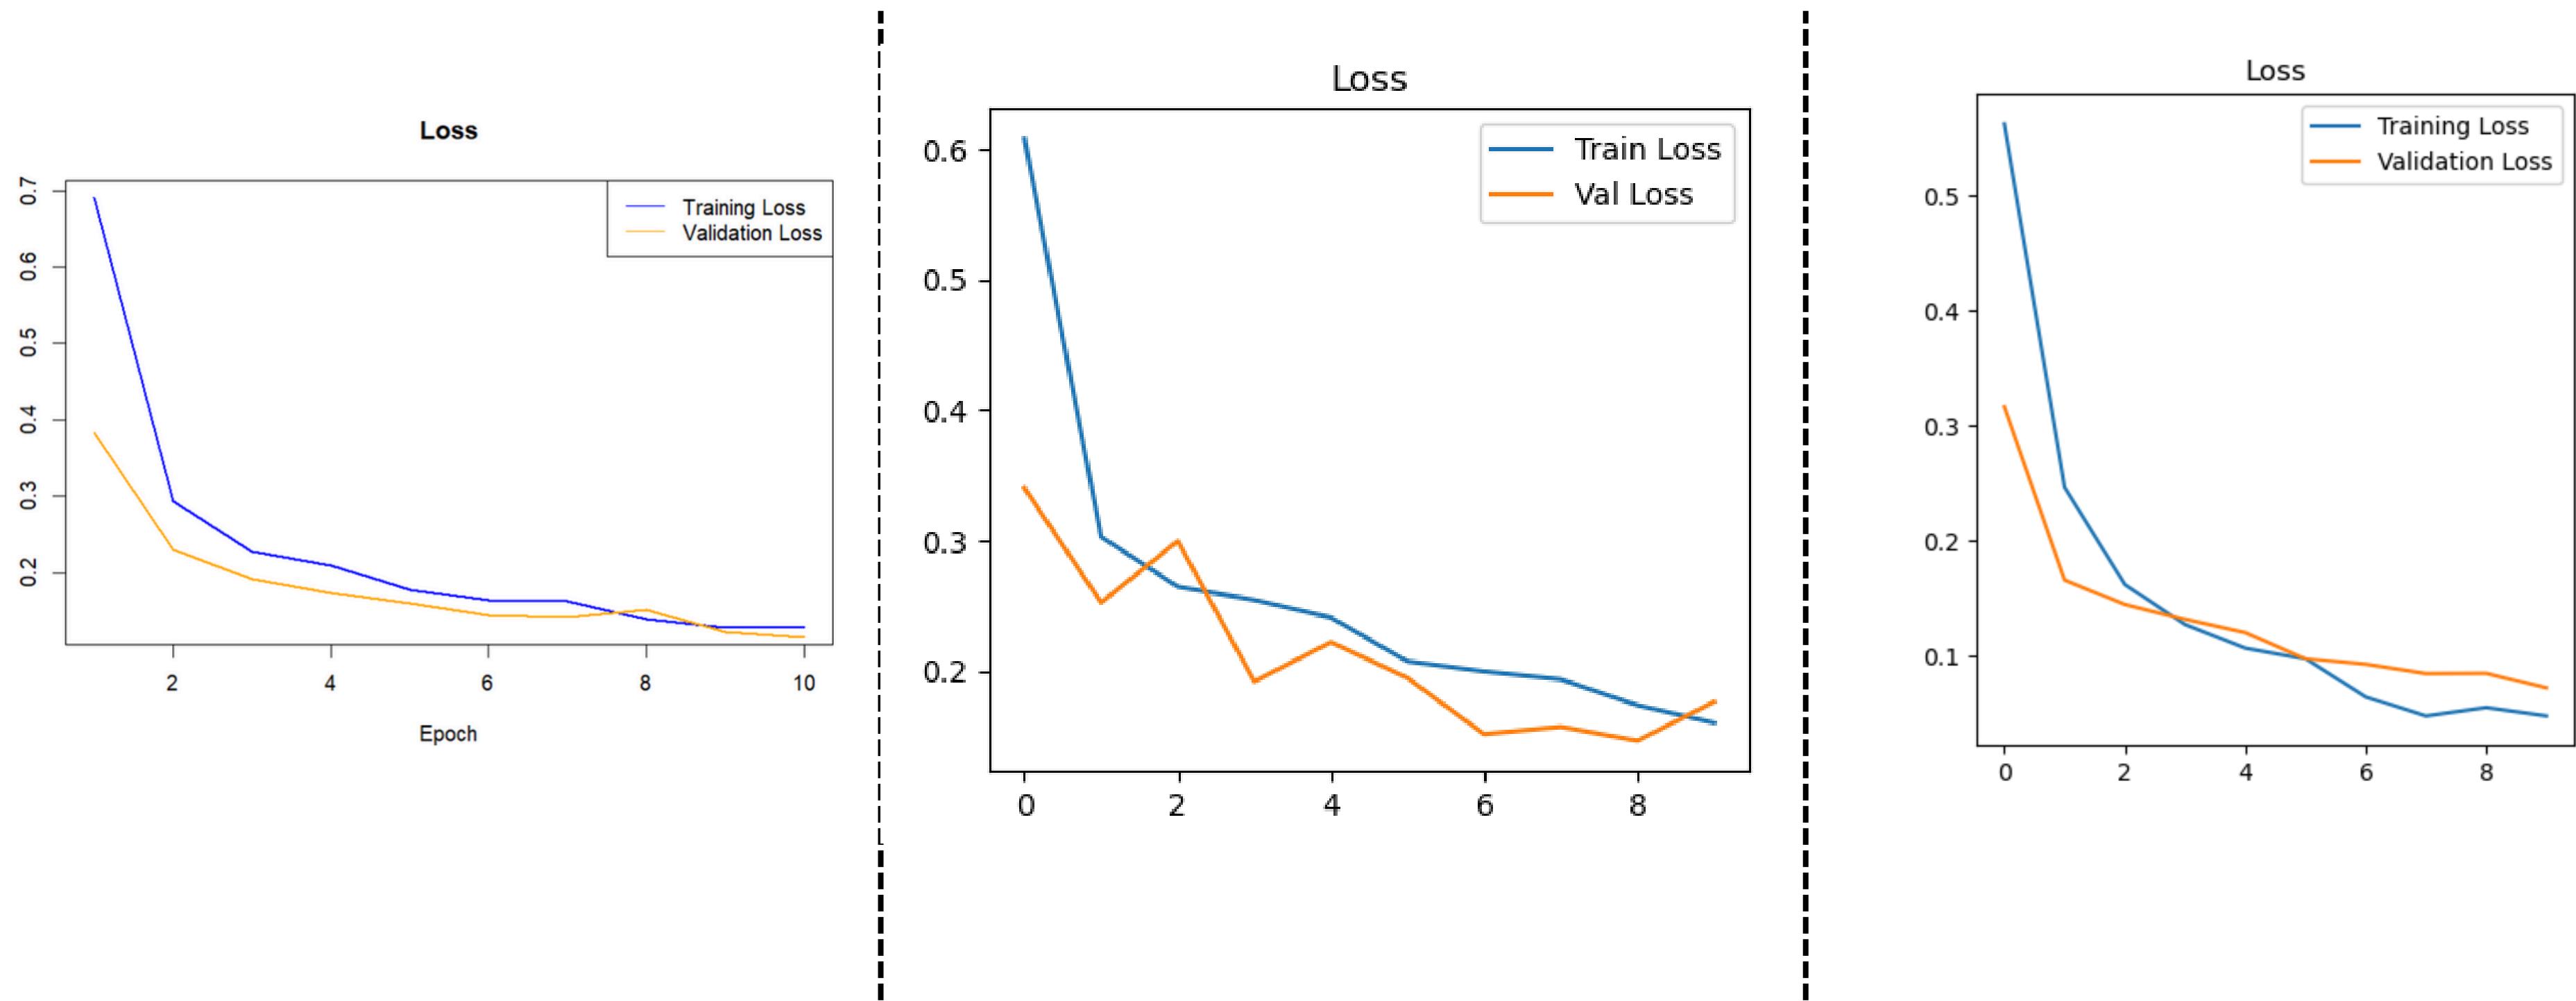
Le modèle d'apprentissage se fait via des boucles explicites manuelles

Plus flexible, plus de code mais plus de maîtrise pour la recherche ou tester de nouvelles architectures

Validation du modèle



Validation du modèle



Temps d'entraînement



16 minutes

Performance dépendantes de reticulate

R Studio utilise par défaut un seul cœur



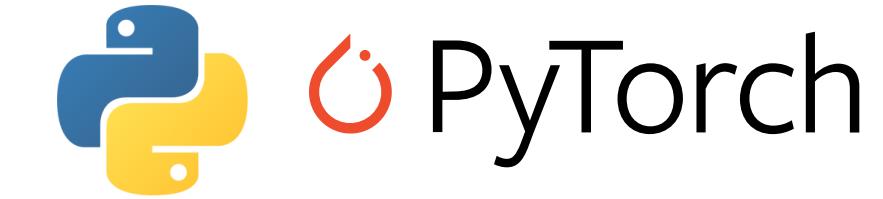
Très rapide : 1 min 30

Model.fit() → Boucle d'entraînement compilée (C++/CUDA)

Keras active automatiquement :

- cache (les images sont gardées en mémoire quand possible)
- prefetch (chargement du batch suivant pendant que le GPU travaille)

Les transformations intégrées sont optimisées → s'exécutent dans la pipeline TensorFlow (pas sur python)



Un peu plus long : 8 min

Boucle d'entraînement manuelle en Python → plus de surcoût à chaque batch

Le DataLoader n'est pas automatiquement optimisé → il faut une configuration manuelle

Il y a des possibilités d'optimiser le temps de calcul :

- num_workers
- pin_memory
- cuDNN benchmark

Ressources disponibles



Documentation officielle

+ CRAN

Getting Started with Keras

Overview

Keras is a high-level neural networks API developed with a focus on enabling fast experimentation. Being able to go from idea to result with the least possible delay is key to doing good research. Keras has the following key features:

- Allows the same code to run on CPU or on GPU, seamlessly.
- User-friendly API which makes it easy to quickly prototype deep learning models.
- Built-in support for convolutional networks (for computer vision), recurrent networks (for sequence processing), and any combination of both.
- Supports arbitrary network architectures: multi-input or multi-output models, layer sharing, model sharing, etc. This means that Keras is appropriate for building essentially any deep learning model, from a memory network to a neural Turing machine.

This website provides documentation for the R interface to Keras. See the main Keras website at <https://keras.io> for additional information on the project.

Nombreux guides

The screenshot shows the TensorFlow for R documentation page. At the top, there's a navigation bar with links for 'TensorFlow for R', 'Install', 'Tutorials', 'Guides', 'Examples', and 'Reference'. Below the navigation, there's a sidebar with a tree view of 'Guides' including 'Tensorflow Basics', 'Tensors', 'Variables', 'Automatic differentiation', 'Graphs and functions', 'Keras', and 'Guide to Keras Basics'. The main content area displays a 'Guide to Keras Basics' section with text about Keras being a high-level API for building and training deep learning models, its advantages (User friendly, Modular and composable, Easy to extend), and how to get started with 'library(keras)'. There are also sections for 'Import keras' and 'Build a simple model'.

StackOverFlow

'keras' + 'r' → > 600 questions

'keras' + 'CNN' + 'r' → > 30 questions



Documentation officielle

The screenshot shows the Keras 3 API documentation page. At the top, there's a navigation bar with links for 'GET STARTED', 'GUIDES', 'API', 'EXAMPLES', 'KERAS TUNER', 'KERAS RS', 'KERAS HUB', and 'SEARCH'. The main content area is titled 'Keras 3 API documentation' and contains two main sections: 'Models API' and 'Layers API'. The 'Models API' section lists items like 'The Model class', 'The Sequential class', 'Model training APIs', and 'Saving & serialization'. The 'Layers API' section lists items like 'The base Layer class', 'Layer activations', 'Layer weight initializers', 'Layer weight regularizers', 'Core layers', 'Convolution layers', 'Pooling layers', 'Recurrent layers', 'Preprocessing layers', 'Normalization layers', and 'Regularization layers'.

Immense communauté

StackOverFlow

'keras' + 'python' → > 26 000 questions

'keras' + 'CNN' + 'python' → > 2 000 questions



Documentation officielle

The screenshot shows the PyTorch documentation page. At the top, there's a navigation bar with links for 'v2.9.1 (stable)', 'Home', 'Install PyTorch', 'User Guide', 'Reference API', 'Developer Notes', 'Community', and 'Tutorials'. The main content area is titled 'PyTorch documentation' and contains sections for 'PyTorch Libraries' (listing torchao, torchrect, torchfit, TorchCodeC, torchvision, ExecuTorch, PyTorch on XLA Devices), 'Install PyTorch', 'User Guide', and 'Pytorch Overview'. On the right side, there's a sidebar with links for 'On this page', 'Indices and tables', 'Edit on GitHub', 'Show Source', and 'Ask AI'.

Immense communauté

StackOverFlow

'Pytorch' → ≈ 24 000 questions

'Pytorch' + 'CNN' → > 1 400 questions

Conclusion



Installation compliquée
Gestion des données lourde
Risque de confusion entre keras et keras3
Ressources importante
Syntaxe simple

+ rapide
- performant



Syntaxe simple
Optimisation déjà faite
Ressources importante

+ rapide
+ performant



Syntaxe plus complète
Difficile d'optimiser les calculs
Meilleur contrôle

- rapide
+ performant



Sources



- R Interface to Keras - <https://keras.rstudio.com/>
- CRAN Getting started with keras <https://cran.r-project.org/web/packages/keras/vignettes/>
- Karlijn Willems, « keras : Deep Learning in R », Juin 2017



Merci

Avez-vous des questions ?